Natural Language Processing & Word Embeddings

Quiz, 10 questions

10/10 points (100%)

✔ **Congratulations! You passed!**

Next Item

✔ 1 / 1
points

1.
Suppose you learn a word embedding for a vocabulary of 10000 words.
Then the embedding vectors should be 10000 dimensional, so as to capture
the full range of variation and meaning in those words.

◯ True

◯ False

Correct

✔ 1 / 1
points

2.
What is t-SNE?

◯ A linear transformation that allows us to solve analogies on word
vectors

◯ A non-linear dimensionality reduction technique

Correct

◯ A supervised learning algorithm for learning word embeddings

◯ An open-source sequence modeling library

Natural Language Processing & Word Embeddings

10/10 points (100%)

Quiz, 10 questions

✓   1 / 1
     points

3.

Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word "ecstatic" does not appear in your small training set, your RNN might reasonably be expected to recognize "I'm ecstatic" as deserving a label $y = 1$.

○ True

Correct

○ False

✓   1 / 1
     points

4.

Which of these equations do you think should hold for a good word embedding? (Check all that apply)

☐ $e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$

Correct

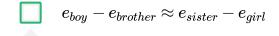☐ $e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$

Natural Language Processing & Word Embeddings

Un-selected is correct

Quiz, 10 questions

10/10 points (100%)

☐ $e_{boy} - e_{brother} \approx e_{girl} - e_{sister}$

Correct

☐ $e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$

Un-selected is correct

✔ 1 / 1
points

5.

Let E be an embedding matrix, and let e_{1234} be a one-hot vector corresponding to word 1234. Then to get the embedding of word 1234, why don't we call $E * e_{1234}$ in Python?

○ It is computationally wasteful.

Correct

○ The correct formula is $E^T * e_{1234}$.

○ This doesn't handle unknown words (<UNK>).

○ None of the above: Calling the Python snippet as described above is fine.

✔ 1 / 1
points

6.

When learning word embeddings, we create an artificial task of estimating $P(target \mid context)$. It is okay if we do poorly on this artificial prediction

task, since the more important by-product of this task is that we learn a useful set

of word embeddings.

◯ True

Correct

◯ False

✔ 1 / 1 points

7.

In the word2vec algorithm, you estimate $P(t \mid c)$, where t is the target word and c is a context word. How are t and c chosen from the training set? Pick the best answer.

◯ c is a sequence of several words immediately before t.

◯ c is the sequence of all the words in the sentence before t.

◯ c is the one word that comes immediately before t.

◉ c and t are chosen to be nearby words.

Correct

✔ 1 / 1 points

8.

Natural Language Processing & Word Embeddings

Quiz, 10 questions

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

Which of these statements are correct? Check all that apply.

- ☑ θ_t and e_c are both 500 dimensional vectors.

  Correct

- ☐ θ_t and e_c are both 10000 dimensional vectors.

  Un-selected is correct

- ☑ θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.

  Correct

- ☐ After training, we should expect θ_t to be very close to e_c when t and c are the same word.

  Un-selected is correct

✔ 1 / 1
points

9.

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings.The GloVe model minimizes this objective:

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - log X_{ij})^2$$

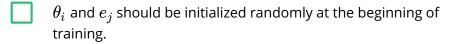Which of these statements are correct? Check all that apply.

- ☐ θ_i and e_j should be initialized to 0 at the beginning of training.

Un-selected is correct

Natural Language Processing & Word Embeddings

10/10 points (100%)

Quiz, 10 questions

☐ θ_i and e_j should be initialized randomly at the beginning of training.

▲

Correct

☐ X_{ij} is the number of times word i appears in the context of word j.

▲

Correct

☐   The weighting function $f(.)$ must satisfy $f(0) = 0$.

▲

Correct

The weighting function helps prevent learning only from extremely common word pairs. It is not necessary that it satisfies this function.

✔   1 / 1
     points

10.

You have trained word embeddings using a text dataset of m_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of m_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstance would you expect the word embeddings to be helpful?

◯ $m_1 \gg m_2$

▲

Correct

◯ $m_1 \ll m_2$